Boosting With the L 2 Loss

2 authors:

Peter Bühlmann
ETH Zurich
282 PUBLICATIONS   **13,720** CITATIONS

SEE PROFILE

B. Yu
University of California, Berkeley
179 PUBLICATIONS   **13,357** CITATIONS

SEE PROFILE

Some of the authors of this publication are also working on these related projects:

Variable Length Markov Chains: R implementation in VLMC View project

Characterization and Greedy Learning of Interventional Markov Equivalence Classes of Directed Acyclic Graphs View project

Boosting with the L_2-Loss: Regression and Classification

Peter Bühlmann
ETH Zürich

Bin Yu
University of California, Berkeley

June 2002
(Revised Version)

Abstract

This paper investigates a computationally simple variant of boosting, L_2Boost, which is constructed from a functional gradient descent algorithm with the L_2-loss function. As other boosting algorithms, L_2Boost uses many times in an iterative fashion a pre-chosen fitting method, called the learner. Based on the explicit expression of refitting of residuals of L_2Boost, the case with (symmetric) linear learners is studied in detail in both regression and classification. In particular, with the boosting iteration m working as the smoothing or regularization parameter, a new exponential bias-variance trade off is found with the variance (complexity) term increasing very slowly as m tends to infinity. When the learner is a smoothing spline, an optimal rate of convergence result holds for both regression and classification and the boosted smoothing spline even adapts to higher order, unknown smoothness. Moreover, a simple expansion of a (smoothed) 0-1 loss function is derived to reveal the importance of the decision boundary, bias reduction, and impossibility of an additive bias-variance decomposition in classification. Finally, simulation and real data set results are obtained to demonstrate the attractiveness of L_2Boost. In particular, we demonstrate that L_2Boosting with a novel component-wise cubic smoothing spline is both practical and effective in the presence of high-dimensional predictors.

1 Introduction

Boosting is one of the most successful and practical methods that recently come from the machine learning community. Since its inception in 1990 (Schapire, 1990; Freund, 1995; Freund and Schapire, 1996), it has been tried on an amazing array of data sets. The improved performance through boosting of a fitting method, called the learner, has been impressive, and it seems to be associated with boosting's resistance to overfitting. The burning question is why.

The rationale behind boosting separates itself from the traditional statistical procedures. It starts with a sensible estimator or classifier, the learner, and seeks its improvements iteratively based on its performance on the training data set. The possibility of this boosting procedure comes with the availability of large data sets where one can easily set aside part of it as the test set (or use cross validation based on random splits). It seemingly bypasses the need to get a model for the data and the pursuit of the optimal solution under this model as the common practice in traditional statistics. For large data set problems with high-dimensional predictors, a good model for the problem is hard to come by, but a sensible procedure is not. And this may explain the empirical success of boosting on large, high-dimensional data sets. After much work on bounding the test set

1

error (generalization error) of a boosted procedure via the VC dimensions and the distribution of so-called margins (Schapire et al., 1998), some recent developments on boosting have been on the gradient-descent view of boosting. They are the results of efforts of many researchers (Breiman, 1999; Mason et al., 1999; Friedman et al., 2000; Collins et al., 2000). This gradient descent view connects boosting to the more common optimization view of statistical inference, and its most obvious consequence has been the emergence of many variants of the original AdaBoost, under various loss or objective functions (Mason et al., 1999; Friedman et al., 2000; Friedman, 2001). Even though a satisfactory explanation on why boosting works does not follow directly from this gradient descent view, some of the new boosting variants are more easily accessible for analysis. In this paper, we take advantage of this new analytic possibility on L_2-boosting procedures to build our case for understanding boosting both in regression and two-class classification. It is worth pointing out that L_2Boost is studied here also as a procedure yielding competitive statistical results in regression and classification, in addition to its computational simplicity and analytical tractability.

After a brief overview of boosting from the gradient descent point of view in Section 2, Section 3 deals with the case of (symmetric) linear learners in regression, building on the known fact that L_2Boost is a stagewise refitting of the residuals (cf. Friedman, 2001). We derive two main rigorous results:

(i) With the boosting iteration m working as the smoothing or regularization parameter, a new exponential bias-variance trade off is found. When the iteration m increases by 1, one more term is added in the fitted procedure, but due to the dependence of this new term on the previous terms, the "complexity" of the fitted procedure is not increased by a constant amount as we got used to in linear regression, but an exponentially diminishing amount as m gets large. At the iteration limit, the complexity or variance term is bounded by the noise variance in the regression model.

(ii) When the learner is a smoothing spline, L_2Boost achieves the optimal rate of convergence for one-dimensional function estimation. Moreover, this boosted smoothing spline adapts to higher order, *unknown* smoothness.

Item (i) partially explains the "overfitting-resistance" mystery of boosting. The phenomenon is radically different from the well-known algebraic bias-variance trade-off in nonparametric regression. Item (ii) shows an interesting result about boosting in adaptive estimation: even when smoothness is unknown, L_2Boost achieves the optimal (minimax) rate of convergence.

Section 4 proposes L_2Boost with a novel component-wise smoothing spline learner as a very effective procedure to carry out boosting for high dimensional regression problems with continuous predictors. It is shown to outperform L_2Boost with stumps (tree with two terminal nodes) and other more traditional competitors, particularly when the predictor space is very high-dimensional.

Section 5 deals with classification, first with the two class problem and then the multiclass problem using the "one against all" approach. The optimality in item (ii) above also holds for classification: L_2Boost achieves the optimal (minimax) rate of convergence to the Bayes risk over an appropriate smoothness function class, the risk of the best among all classification procedures. Furthermore in Section 6, we approximate the 0-1 loss function via a smoothed version to show that the test set (generalization) error of any procedure is approximately, in addition to the Bayes risk, a sum of tapered moments. As

a consequence of this approximation, we get more insight into why bias plays a bigger role in 0-1 loss classification than in L_2-regression, why there is even more "resistance against overfitting" in classification than regression and why previous attempts were not successful at decomposing the test set (generalization) error into additive bias and variance terms (cf. Geman et al. 1992; Breiman, 1998, and references therein).

In Section 7, we support the theory and explanations for classification by simulated and real data sets which demonstrate the attractiveness of L_2Boost. Finally, Section 8 contains a discussion on the role of the learner and a summary of the paper.

2   Boosting: stagewise functional gradient descent

The boosting algorithms can be seen as functional gradient descent techniques. The task is to estimate the function $F : \mathbb{R}^d \to \mathbb{R}$, minimizing an expected cost

$$\mathbb{E}[C(Y, F(X))], \ C(\cdot, \cdot) : \mathbb{R} \times \mathbb{R} \to \mathbb{R}^+ \tag{1}$$

based on data (Y_i, X_i) $(i = 1, \ldots n)$. We consider here both cases where the univariate response Y is continuous (regression problem) or discrete (classification problem), since boosting is potentially useful in both cases; X denotes here a d-dimensional predictor variable. The cost function $C(\cdot, \cdot)$ is assumed to be smooth and convex in the second argument to ensure that the gradient method works well. The most prominent examples are:

$$C(y, f) = \exp(yf) \text{ with } y \in \{-1, 1\}: \text{ AdaBoost cost function,}$$
$$C(y, f) = \log_2(1 + \exp(-2yf)) \text{ with } y \in \{-1, 1\}: \text{ LogitBoost cost function,}$$
$$C(y, f) = (y - f)^2/2 \text{ with } y \in \mathbb{R} \text{ or } \in \{-1, 1\}: L_2\text{Boost cost function.} \tag{2}$$

The population minimizers of (1) are then (cf. Friedman et al., 2000)

$$F(x) = \frac{1}{2} \log(\frac{\mathbb{P}[Y = 1 | X = x]}{\mathbb{P}[Y = -1 | X = x]}) \text{ for AdaBoost and LogitBoost cost,}$$
$$F(x) = \mathbb{E}[Y | X = x] \text{ for } L_2\text{Boost cost.} \tag{3}$$

Estimation of such an $F(\cdot)$ from data can be done via a constrained minimization of the empirical risk

$$n^{-1} \sum_{i=1}^{n} C(Y_i, F(X_i)). \tag{4}$$

by applying functional gradient descent. This gradient descent view has been recognized and refined by various authors including Breiman (1999), Mason et al. (1999), Friedman et al. (2000), Friedman (2001). In summary, the minimizer of (4) is imposed to satisfy a "smoothness" (or "regularization") constraint in terms of an additive expansion of ("simple") learners (fitted functions)

$$h(x, \hat{\theta}), \ x \in \mathbb{R}^d,$$

where $\hat{\theta}$ is an estimated finite or infinite-dimensional parameter. For example, the learner $h(\cdot, \hat{\theta})$ could be a decision tree where $\hat{\theta}$ describes the axis to be split, the split points and

the fitted values for every terminal node (the constants in the piecewise constant fitted function) How to fit $h(x, \theta)$ from data is part of the learner and can be done according to a basis algorithm. For example, least squares fitting yields

$$\hat{\theta}_{U,X} = \text{argmin}_\theta \sum_{i=1}^{n} (U_i - h(X_i; \theta))^2,$$

for some data $(U, X) = \{(U_i, X_i); i = 1, \ldots, n\}$. The general description of functional gradient descent is as follows (cf. Friedman, 2001).

Generic functional gradient descent

Step 1 (initialization). Given data $\{(Y_i, X_i); i = 1, \ldots, n\}$, fit a real-valued, (initial) learner

$$\hat{F}_0(x) = h(x; \hat{\theta}_{Y,X}).$$

When using least squares, $\hat{\theta}_{Y,X} = \text{argmin}_\theta \sum_{i=1}^{n} (Y_i - h(X_i; \theta))^2$. Set $m = 0$.

Step 2 (projecting gradient to learner). Compute the negative gradient vector

$$U_i = -\frac{\partial C(Y_i, F)}{\partial F}\big|_{F = \hat{F}_m(X_i)}, \ i = 1, \ldots, n,$$

evaluated at the current $\hat{F}_m(\cdot)$. Then, fit the real-valued learner to the gradient vector

$$\hat{f}_{m+1}(x) = h(x, \hat{\theta}_{U,X}).$$

When using least squares, $\hat{\theta}_{U,X} = \text{argmin}_\theta \sum_{i=1}^{n} (U_i - h(X_i; \theta))^2$.

Step 3 (line search). Do one-dimensional numerical search for the best step-size

$$\hat{w}_{m+1} = \text{argmin}_w \sum_{i=1}^{n} C(Y_i, \hat{F}_m(X_i) + w_{m+1}\hat{f}_{m+1}(X_i)).$$

Update,

$$\hat{F}_{m+1}(\cdot) = \hat{F}_m(\cdot) + \hat{w}_{m+1}\hat{f}_{m+1}(\cdot).$$

Step 4 (iteration). Increase m by one and repeat Steps 2 and 3.

The learner $h(x, \hat{\theta}_{U,X})$ in Step 2 can be viewed as an estimate of $\mathbb{E}[U_i | X = x]$ and takes values in $\mathbb{R}$, even in case of a classification problem with Y_i in a finite set. We call $\hat{F}_m(\cdot)$ the AdaBoost-, LogitBoost- or L_2Boost-estimate, according to the implementing cost function in (2).

L_2Boost has a simple structure: the negative gradient in Step 2 is the classical residual vector and the line search in Step 3 is trivial.

L_2Boost algorithm

Step 1 (initialization). As in Step 1 of generic functional gradient descent, using a least squares fit.

Step 2. Compute residuals $U_i = Y_i - \hat{F}_m(X_i)$ $(i = 1, \dots, n)$ and fit the real-valued learner to the current residuals by least squares as in Step 2 of the generic functional gradient descent; the fit is denoted by $\hat{f}_{m+1}(\cdot)$.
Update

$$\hat{F}_{m+1}(\cdot) = \hat{F}_m(\cdot) + \hat{f}_{m+1}(\cdot).$$

Step 3 (iteration). Increase iteration index m by one and repeat Step 2.

L_2Boosting is thus nothing else than repeated least squares fitting of residuals (cf. Friedman, 2001). With $m = 1$ (one boosting step), it has already been proposed by Tukey (1977) under the name "twicing".

For a continuous $Y \in \mathbb{R}$, a regression estimate for $\mathbb{E}[Y|X = x]$ is directly given by the L_2Boost-estimate $\hat{F}_m(\cdot)$. For a two-class problem with $Y \in \{-1, 1\}$, a classifier under equal misclassification costs is given by

$$sign(\hat{F}_m(x)) \tag{5}$$

since $\mathbb{E}[Y|X = x] = \mathbb{P}[Y = 1|X = x] - \mathbb{P}[Y = -1|X = x]$. AdaBoost- and LogitBoost-estimates aim to estimate

$$F(x) = \frac{1}{2} \log \left(\frac{\mathbb{P}[Y = 1|X = x]}{\mathbb{P}[Y = -1|X = x]} \right).$$

Hence, an appropriate classifier is again given by (5).

Mason et al. (1999) and Collins et al. (2000) describe when boosting-type algorithms, i.e. functional gradient descent, converge numerically. This tells us that, under certain conditions, the test set (generalization) error for boosting eventually stabilizes. But it doesn't imply that the eventually stable solution is the best, or that overfitting could happen long before reaching convergence. Indeed, we will show in Section 3 that L_2Boost with "contracting" linear learners converges to the fully saturated model, i.e. $\hat{F}_\infty(X_i) = Y_i$ for all $i = 1, \dots, n$, fitting the data exactly.

Obviously, L_2Boost and other functional gradient descent methods depend on the choice of the learner. As the boosting iteration runs, the boosted procedure has more terms and hence becomes more complex. It is intuitively clear that L_2boosting is not worthwhile if the learner is already complex, say fitting many parameters, so that every boosting iteration contributes to additional overfitting. We make this rigorous in Section 3 for linear learners. Thus, the learner in boosting should be "simple": it typically involves only few parameters and has low variance relative to bias. We say that such a learner is weak, an informal terminology from machine learning. Weakness of a learner does depend on the signal to noise ratio of the data: if the noise level is low, it is well known that a statistical method has less a tendency to overfit. Of course, there are many possibilities for choosing a weak learner: examples include stumps which are trees with two terminal nodes only, smoothing methods with large amount of smoothing, or shrinking a learner with a small shrinkage factor. Some illustrations are given in Sections 3.2.2, 4.2, 7 and 8.

3 L_2Boosting with linear learners in regression

The nature of stagewise fitting is responsible to a large extent for boosting's resistance to overfitting. The same view has been expressed in Buja's (2000) discussion of the Friedman et al. (2000) paper. He made amply clear there that this stagewise fitting had gotten a bad reputation among statisticians and did not get the attention it deserved. The success of boosting definitely serves as an eye-opener for us to take a fresh look at stagewise fitting.

3.1 Theory

Consider the regression model

$$Y_i = f(x_i) + \varepsilon_i, \ i = 1, \ldots, n,$$
$$\varepsilon_1, \ldots, \varepsilon_n \text{ i.i.d. with } \mathbb{E}[\varepsilon_i] = 0, \ \text{Var}(\varepsilon_i) = \sigma^2, \tag{6}$$

where $f(\cdot)$ is a real-valued, typically nonlinear function, and the predictors $x_i \in \mathbb{R}^d$ are deterministic (e.g. conditioning on the design).

We can represent a learner, evaluated at the predictors $x_1, \ldots, x_n$ as an operator $\mathcal{S} : \mathbb{R}^n \to \mathbb{R}^n$, mapping the responses $Y_1, \ldots, Y_n$ to some fitted values in $\mathbb{R}^n$. The predictors $x_1, \ldots, x_n$ are absorbed in the operator notation $\mathcal{S}$. In the sequel, we often use the notation Y for the vector $(Y_1, \ldots, Y_n)^T$, F_j for the vector $(F_j(x_1), \ldots, F_j(x_n))^T$ and analogously for f_j; it should always be clear from the context whether we mean a single variable Y or function $F_j(\cdot)$, or the vectors as above.

Proposition 1. *The L_2Boost estimate in iteration m can be represented as:*

$$\hat{F}_m = \sum_{j=0}^{m} \mathcal{S}(I - \mathcal{S})^j Y = (I - (I - \mathcal{S})^{m+1})Y.$$

A proof is given in the Appendix. We define the boosting operator $\mathcal{B}_m : \mathbb{R}^n \to \mathbb{R}^n$ by

$$\mathcal{B}_m Y = \hat{F}_m.$$

Proposition 1 describes an explicit relationship between the boosting operator and the learner $\mathcal{S}$. We exploit this in the sequel.

We focus now on linear learners $\mathcal{S}$. Examples include least squares fitting in linear models, more general projectors to a given class of basis functions such as regression splines, or smoothing operators such as kernel and smoothing spline estimators.

Proposition 2. *Consider a linear learner $\mathcal{S}$ with eigenvalues $\{\lambda_k; \ k = 1, \ldots, n\}$, based on deterministic predictors $x_1, \ldots, x_n$. Then, the eigenvalues of the L_2Boost operator $\mathcal{B}_m$ are $\{(1 - (1 - \lambda_k)^{m+1}; \ k = 1, \ldots, n\}$.*

Proof: This is a direct consequence of Proposition 1. □

Our analysis will become even more transparent when specialized to the case where $\mathcal{S} = \mathcal{S}^T$ is symmetric. An important example is the smoothing spline operator (see Wahba, 1990; Hastie and Tibshirani, 1990) which is a more data-adaptive smoothing technique than say kernel with a global bandwidth. All eigenvalues of $\mathcal{S}$ are then real and $\mathcal{S}$ as well as $\mathcal{B}_m$ can be diagonalized with an orthonormal transform,

$$\mathcal{B}_m = U D_m U^T, \ D_m = \text{diag}(1 - (1 - \lambda_k)^{m+1}),$$
$$k\text{th column-vector of } U \text{ being the } k\text{th eigenvector of } \mathcal{S} \text{ to the eigenvalue } \lambda_k. \tag{7}$$

The matrix U is orthonormal, satisfying $UU^T = U^T U = I$.

We are now able to analyze a relevant generalization measure in this setting, the (expected) mean squared error

$$MSE = n^{-1} \sum_{i=1}^{n} \mathbb{E}[(\hat{F}_m(x_i) - f(x_i))^2], \tag{8}$$

which averages over the observed predictors. Note that if the design is stochastic with a probability distribution, the MSE measure above is asymptotically equivalent (as $n \to \infty$) to the prediction (generalization) error $\mathbb{E}[(\hat{F}_m(X) - f(X))^2]$, where X is a new test observation from the design generating distribution but independent from the training set and the expectation is over the training and test set. We show in Figure 1 the difference between the two measures for a finite sample case.

Proposition 3. *Consider a linear, symmetric learner $\mathcal{S} = \mathcal{S}^T$ with eigenvalues $\{\lambda_k; \ k = 1, \ldots, n\}$ and eigenvectors building the columns of the orthonormal matrix U. Assume data being generated from the model (6) and denote by $f = (f(x_1), \ldots, f(x_n))^T$ the vector of the true regression function $f(\cdot)$ evaluated at x_i's. Then, the squared bias, variance and averaged mean squared error for L_2Boost are*

$$bias^2(m, \mathcal{S}; f) = n^{-1} \sum_{i=1}^{n} (\mathbb{E}[\hat{F}_m(x_i)] - f(x_i))^2 = n^{-1} f^T U \, diag((1 - \lambda_k)^{2m+2}) U^T f,$$

$$variance(m, \mathcal{S}; \sigma^2) = n^{-1} \sum_{i=1}^{n} Var(\hat{F}_m(x_i)) = \sigma^2 n^{-1} \sum_{k=1}^{n} (1 - (1 - \lambda_k)^{m+1})^2,$$

$$MSE(m, \mathcal{S}; f, \sigma^2) = bias^2(m, \mathcal{S}; f) + variance(m, \mathcal{S}; \sigma^2).$$

A proof is given in the Appendix. Proposition 3 describes an exact result for the MSE in terms of the chosen L_2Boost procedure. It is clear that the iteration index m acts as a "smoothing parameter" to control the bias and variance trade-off.

Given the underlying problem (i.e. f and σ^2) and given a learner $\mathcal{S}$ (implying U and the set of eigenvalues), we analyze the bias-variance trade-off as a function of boosting iterations. For that purpose, we assume that all eigenvalues satisfy $0 < \lambda_k \leq 1$. An important example for such a linear learner are cubic smoothing splines which have two eigenvalues equal to one and all others strictly between zero and one: this will be treated even in more detail in Section 3.2.

Theorem 1. *Under the assumptions in Proposition 3 with $0 < \lambda_k \leq 1$, $k = 1, \ldots, n$,*

(1) *$bias^2(m; \mathcal{S}; f)$ decays exponentially fast with increasing m,*
    $variance(m, \mathcal{S}; \sigma^2)$ exhibits exponentially small increase with increasing m, and $\lim_{m \to \infty} MSE(m, \mathcal{S}; f, \sigma^2) = \sigma^2$.

(2) *Moreover, let $\mu = U^T f = (\mu_1, ..., \mu_n)^T$ be the function vector in the linear space spanned by column vectors of U (μ represents f in the coordinate-system of the eigenvectors of $\mathcal{S}$).*
    (i) If $\mu_k^2 / \sigma^2 > 1/(1 - \lambda_k)^2 - 1$ for all k with $\lambda_k < 1$, then boosting improves the MSE over the linear learner $\mathcal{S}$.
    (ii) If $\lambda_k < 1$ for at least one $k \in \{1, \ldots, n\}$ ($\mathcal{S}$ is not the identity operator I), there is an m, such that at the mth iteration, the boosting MSE is strictly lower than σ^2.

7

Figure 1: Mean squared prediction error $\mathbb{E}[(Y - \hat{f}(X))^2]$ (solid line) and MSE criterion from (8) (dotted line) from 100 simulations of model (11) with design uniformly distributed on $[-1/2, 1/2]$, each with $n = 100$. Left: L_2Boost with cubic smoothing spline having df=3, as a function of boosting iterations m. Right: Cubic smoothing spline for various degrees of freedom (various amount of smoothing).

Assertion (1) is a direct consequence of Proposition 3. A proof of assertion (2) is given in the Appendix.

Theorem 1 comes as a surprise: it shows a very interesting bias-variance trade-off that has not been seen in the literature. As the boosting iteration (or smoothing parameter) m varies, both the bias and variance (complexity) term change exponentially with the bias decreasing exponentially fast and the variance increasing with exponentially diminishing terms eventually. This contrasts the standard algebraic trade-off commonly seen in nonparametric estimation. Figure 1 illustrates the difference for a cubic smoothing spline learner (for data from model (11) in Section 3.2.2). The exponential trade-off not only gets very close to the optimal of the MSE of that from the smoothing splines (by varying the smoothing parameter), but also stays really flat afterwards, due to the exponential increase and decrease in the bias and variance terms. Condition $\mu_k^2/\sigma^2 > 1/(1 - \lambda_k)^2 - 1$ in part 2(i) can be interpreted as follows. A large left hand side μ_k/σ^2 mean that f is relatively complex compared to the noise level σ^2. A small right hand side $1/(1 - \lambda_k)^2 - 1$ means that λ_k is small, that is, the learner employs very strong shrinkage or smoothing in the kth eigenvector-direction, or the learner is actually weak in the kth direction. Thus, for boosting to bring improvement, μ_k has to be large relative to the noise level σ^2 and/or λ_k is restricted to be sufficiently small. Hence we see improvements for boosting with weak learners most of the time. The assertion in 2(ii) shows that boosting always beats the unbiased estimator Y as does James-Stein estimator in the space spanned by U.

Boosting until theoretical convergence with $m = \infty$ is typically not a good advice: the MSE with $m = \infty$ is not smaller than the noise level σ^2. The reason is that boosting infinitely often yields the fully saturated model which fits the data exactly. Therefore, one should monitor an estimate of MSE, for example by using a test set or cross-validation.

8

Finally, boosting sometimes does not change the procedure at all.

Corollary 1. *Under the assumptions in Proposition 3 with $\lambda_k \in \{0, 1\}$, $k = 1, \ldots, n$, or equivalently, $\mathcal{S}$ is a linear projection operator, $\mathcal{B}_m \equiv \mathcal{S}$ for $m = 1, 2, \ldots$.*

The phenomenon in Theorem 1 generalizes qualitatively to higher order moments.

Theorem 2. *Under the assumptions in Theorem 1 with $0 < \lambda_k \le 1$, $k = 1, \ldots, n$ and assuming $\mathbb{E}|\varepsilon_1|^p < \infty$ for $p \in \mathbb{N}$,*

$$n^{-1} \sum_{i=1}^{n} \mathbb{E}[(\hat{F}_m(x_i) - f(x_i))^p] = \mathbb{E}[\varepsilon_1^p] + O(\exp(-Cm)) \ (m \to \infty),$$

where $C > 0$ is a constant independent of m (but depending on n and p).

A proof is given in the Appendix. Theorem 2 will be used later to argue that the expected 0-1 loss in classification also exhibits only exponentially small amount of overfitting as boosting iterations $m \to \infty$.

3.2 Smoothing splines as learners

3.2.1 Theory

A special class of symmetric linear learners are the smoothing spline operators when the predictors are one-dimensional. Denote the function class of the νth order smoothness, defined on an interval $[a, b]$, as

$$\mathcal{W}_2^{(\nu)} = \{f : f \ (\nu - 1)\text{-times continuously differentiable and } \int_a^b [f^{(\nu)}(x)]^2 dx < \infty\}. \quad (9)$$

The space $\mathcal{W}_2^{(\nu)}$ is also known as Sobolev space. Let $\mathcal{S}Y = g_r$ be the smoothing spline solution to the penalized least squares problem

$$g_r = g_r(\lambda) = \operatorname{argmin}_{f \in \mathcal{W}_2^{(r)}} \frac{1}{n} \sum_i [Y_i - f(x_i)]^2 + \lambda \int [f^{(r)}(x)]^2 dx \quad (10)$$

Theorem 3. *(Optimality of L_2Boost for smoothing splines). Consider the model in (6) with $x_i \in \mathbb{R}$. Suppose $\mathcal{S}$ is a smoothing spline learner $g_r(\lambda_0)$ of degree r corresponding to a fixed smoothing parameter λ_0. If the true function f is in $\mathcal{W}_2^{(\nu)}$ with $\nu \ge r$ (ν, $r \in \mathbb{N}$), then there is an $m = m(n) = O(n^{2r/(2\nu+1)}) \to \infty$ such that $\hat{F}_{m(n)}$ achieves the optimal minimax rate $n^{-2\nu/(2\nu+1)}$ of the (smoother than degree r) function class $\mathcal{W}_2^{(\nu)}$ in terms of MSE as defined in (8).*

A proof is given in the Appendix. This result states that boosting smoothing splines is minimax optimal and also adapts to higher order smoothness: even if the smoothing spline learner has only degree r, we can adapt to higher order smoothness ν and achieve the optimal MSE rate. Interestingly, any fixed smoothing parameter λ_0 of the smoothing spline learner can be used: from the asymptotic view, this means that the smoothing parameter λ_0 is large, i.e. a smoothing spline learner learner with low variance.

Gu (1987) analyzes twicing ($m = 1$) and shows that twicing can adapt to a higher order smoothness $\nu \le 2r$. With boosting we can adapt to an arbitrarily higher order

9

smoothness since we can refit as many times as we want. For cubic smoothing spline learners with $r = 2$, the optimal rate $n^{-4/5}$ is achieved by $m = O(n^{4/5})$. If the underlying smoothness is say $\nu = 3 > 2 = r$, then the boosted cubic smoothing spline can achieve the optimal rate $n^{-6/7}$ for the smoother class with $m = O(n^{4/7})$. In practice, the data driven "optimal" boosting iteration m is selected either through a fixed test set or via cross validation. We note here that boosting also adapts to lower order smoothness $\nu < r$. But this is also true for smoothing splines (without boosting). Hence, boosting is not offering more for this case.

For a given smoothness ν, both the ordinary smoothing spline (with $r = \nu$) and boosting achieve the optimal rate, but they trade off bias and variance along different regularization paths. As mentioned already, the advantage of the new exponential trade-off in boosting is the flatter near-optimal region for the optimal smoothing parameter (or boosting iteration). An example was shown in Figure 1 with simulated data from the next Section.

3.2.2 Simulation results with cubic smoothing spline as learners

The relevance of the theoretical results above depends on the underlying problem and the sample size. We consider a representative example for the model in (6),

$$f(x) = 0.8x + \sin(6x), \ x \in \mathbb{R}^1,$$
$$\varepsilon_i \sim \mathcal{N}(0, \sigma^2), \ \sigma^2 = 2, \ \text{sample size } n = 100. \tag{11}$$

The learner $\mathcal{S}$ is chosen as a cubic smoothing spline which satisfies linearity, symmetry and the eigenvalue-conditions used in Theorems 1 and 2.

The complexity of $\mathcal{S}$, or the strength of the leaner, is chosen here in terms of the so-called degrees of freedom (df) which equals the trace of $\mathcal{S}$ (Hastie and Tibshirani, 1990). To study the interaction of the learner with the underlying problem, we fix the model as in (11) and a cubic smoothing-spline learner with df=20. To decrease the learner's complexity (or increase the learner's weakness), we use shrinkage (Friedman, 2001) to replace $\mathcal{S}$ by

$$\mathcal{S}_\nu = \nu\mathcal{S}, \ 0 < \nu \leq 1.$$

For $\mathcal{S}$ a smoothing spline estimator, shrinkage with ν small corresponds to a linear operator $\mathcal{S}_\nu$ whose eigenvalues $\{\nu\lambda_k\}_k$ are closer to zero than $\{\lambda_k\}_k$ for the original $\mathcal{S}$. With small ν, we thus get a weaker learner than the original $\mathcal{S}$: shrinkage acts here similarly as changing the degrees of freedom of the original $\mathcal{S}$ to a lower value. We will see its effect in more complex examples in Sections 4.2 and 8. The boosting question becomes whether even a very weak $\mathcal{S}_\nu$, with ν very small, can be boosted with m large to achieve almost optimal performance (defined through numerical search among all the estimators rising from different shrinkage factors and different iterations of boosting). Figure 2 displays MSE from specification (11) with $x_1, \ldots, x_n$ i.i.d. realizations from $\mathcal{N}(0, 1)$, as a function of m and ν. It shows that the boosting question from above has a positive answer for this case. That is, we observe the following from this example:

(1) Boosting with a large number of iterations has the potential to make a very weak learner (with ν very small) almost optimal when compared with the best shrunken learner $\nu_{opt}S$. This is consistent with the asymptotic result in Theorem 3.
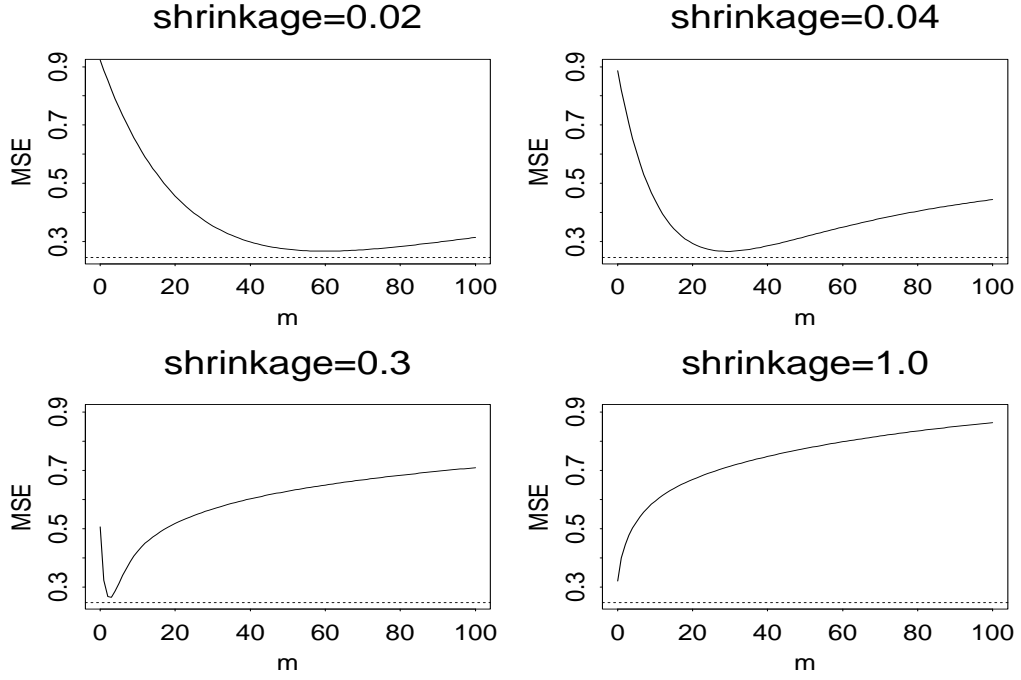
Figure 2: Traces of MSE as a function of boosting iterations m, for four different shrinkage factors. Dotted line represents minimum, achieved with $m = 0, \nu = 0.76$. The data is from model (11) with sample size $n = 100$.

(2) Provided that the learner is sufficiently weak, boosting always improves, as we show in Theorem 1.

(3) When the initial learner $\mathcal{S}$ is too strong, boosting decreases performance due to overfitting. (In our case, the learner with 20 degrees of freedom is too strong or has a too small amount of smoothing.)

(4) Boosting very weak learners is relatively safe, provided that the number of iterations is large: the MSE with ν very low is flat for large number of iterations m.

Statements (2)–(4), numerically found for this example with a linear learner, have also been shown to hold empirically for many data sets and with nonlinear learners in AdaBoost and LogitBoost in classification. Statement (1), dealing with optimality, is asymptotically explained by our Theorem 3. Also the recent work of Jiang (2000) hints at a consistency result for AdaBoost: he shows that in the asymptotic sense and for classification, AdaBoost visits optimal (Bayes-) procedures during the evolution of AdaBoost. However, this asymptotic consistency result doesn't explain optimality or AdaBoost's good finite-sample performance.

All these behaviors happen well before the asymptopia $m = \infty$ or the convergence of the boosting algorithm. For example, when having perfect knowledge of the MSE at every boosting iteration as in our simulation example, a suitable stopping criterion would be the relative difference of MSE between the current and previous iteration. When using the tolerance, say 10^{-4} for this relative difference we would need $m = 761$, when using S_ν with $\nu = 1$, and the resulting MSE at stopping is 1.088. For $\nu = 0.02$, we would need $m = 1691$ and the MSE at stopping is 1.146. This is still far from the stabilizing MSE

sample size n	optimal smoothing spline	optimal L_2Boost	gain
10	$7.787 \cdot 10^{-1}$	$9.968 \cdot 10^{-1}$	-28.0%
25	$3.338 \cdot 10^{-1}$	$3.349 \cdot 10^{-1}$	-0.3%
50	$1.657 \cdot 10^{-1}$	$1.669 \cdot 10^{-1}$	-0.7%
100	$9.332 \cdot 10^{-2}$	$9.050 \cdot 10^{-2}$	0.9%
1000	$1.285 \cdot 10^{-2}$	$1.128 \cdot 10^{-2}$	12.2%

Table 1: Mean squared error for simulated data from (11). Optimal cubic smoothing spline (with best penalty parameter) and optimal L_2Boost with smoothing spline (with best number of boosting iterations). Positive gain, which measures the relative improvement of mean squared error with L_2Boost, indicates an advantage for boosting.

$(m = \infty)$ which is $\sigma^2 = 2$, see Theorem 1. The positive aspect is that we would stop before reaching the extreme overfitting situation. This little illustration describes how slow convergence in boosting could be.

We also demonstrate here empirically the adaptivity of L_2Boost with smoothing splines to higher smoothness, see Theorem 3. We use simulated data from model (11) with $x_1, \ldots, x_n$ i.i.d. realizations from Uniform$([-1/2, 1/2])$, for sample sizes $n = 10$ up to 1000. Table 1 reports the performance of the best cubic smoothing spline (with optimal penalty parameter) in comparison to L_2Boosting a cubic smoothing spline with "fixed" penalty $\lambda_0 = c$, a constant (with optimal number of boosting iterations). We evaluate the mean squared error $\mathbb{E}[(f(X) - \hat{f}(X))^2]$ (X a new test observation) by averaging over 100 simulations from the model (11). We observe in Table 1 the following. For $n = 10$, the smoothing spline learner in L_2Boost is too strong and actually, the best performance is with $m = 0$ (no boosting). In the mid-range with $25 \leq n \leq 100$, the differences between optimal smoothing spline and optimal L_2Boost are negligible. For the large sample size $n = 1000$, we see an advantage of L_2Boost which is consistent with the theory: the underlying regression function in (11) is infinitely smooth and L_2Boost exhibits an adaptation to higher order smoothness.

3.3 Estimating the optimal number of boosting iterations

The number m of boosting iterations is a tuning parameter in L_2- and other boosting methods. Theorem 1 and Figure 1 indicate a certain resistance against overfitting as m gets large and thus the tuning of L_2Boost should not be difficult. Nevertheless, a method is needed to choose the value of m based on data, or to estimate the optimal iteration number m.

A straightforward way which we found to work well is given by a 5-fold cross-validation for estimating the mean squared error $\mathbb{E}[(Y - \hat{F}_m(X))^2]$ in regression or the misclassification error $\mathbb{P}[sign(\hat{F}_m(X)) \neq Y]$ in classification. An estimate $\hat{m}$ for the number of boosting iterations is then given as the minimizer of such a cross-validated error.

4 L_2Boosting for regression in high dimensions

When the dimension d of the predictor space is large, the learner $\mathcal{S}$ is typically nonlinear. In very high dimensions, it becomes almost a necessity to use a learner which is doing some sort of variable selection. One of the most prominent examples are trees.

4.1 Component-wise smoothing spline as the learner

As an alternative to tree learners with two terminal nodes (stumps), we propose here component-wise smoothing splines. A component-wise smoothing spline is defined as a smoothing spline with *one selected* explanatory variable $x_{\hat{\iota}}$ ($\hat{\iota} \in \{1, \ldots, d\}$), where

$$\hat{\iota} = \operatorname{argmin}_{\iota} \sum_{i=1}^{n} (Y_i - \hat{g}_{\iota}(x_{i,\iota}))^2,$$

where $\hat{g}_{\iota}$ is the smoothing spline as defined in (10) using the predictor x_{ι}. Thus, the component-wise smoothing spline learner is given by the function

$$\hat{g}_{\hat{\iota}} : x \mapsto \hat{g}_{\hat{\iota}}(x_{\hat{\iota}}), \ x \in \mathbb{R}^d.$$

Boosting stumps and component-wise smoothing splines yields an additive model whose terms are fitted in a stagewise fashion. The reason being that an additive combination of a stump or a component-wise smoothing spline $\hat{F}_0 + \sum_{m=1}^{M} \hat{f}_m$, with $\hat{f}_m(x)$ depending functionally only on $x_{\hat{\iota}}$ for some component $\hat{\iota} \in \{1, \ldots, d\}$, can be re-expressed as an additive function $\sum_{j=1}^{d} \hat{m}_j(x_j)$, $x \in \mathbb{R}^d$. The estimated functions $\hat{m}_j(\cdot)$ when using boosting are fitted in stagewise fashion and different from the backfitting estimates in additive models (cf. Hastie and Tibshirani, 1990). Boosting stumps or component-wise smoothing splines is particularly attractive when aiming to fit an additive model in very high dimensions with d larger or of the order of sample size n. Boosting has then much greater flexibility to add complexity, in a stagewise fashion, to certain components $j \in \{1, \ldots, d\}$ and may even drop some of the variables (components). We will give examples in Section 4.2 that when the dimension d is large, boosting outperforms the alternative classical additive modeling with variable selection, using backfitting.

4.2 Numerical results

We first consider the often analyzed data set of ozone concentration in the Los Angeles basin which has been also considered in Breiman (1998) in connection with boosting. The dimension of the predictor space is $d = 8$ and sample size is $n = 330$. We compare here L_2Boosting with classical additive models using backfitting and with MARS. L_2Boost is used with stumps and with component-wise cubic smoothing splines having 5 degrees of freedom (cf. Hastie and Tibshirani, 1990) and the number of iterations is estimated by 5-fold cross-validation as described in section 3.3; additive model-backfitting is used with the default smoothing splines in S-Plus; MARS is run by using the default parameters as implemented in S-plus, library(mda). We estimate mean squared prediction error $\mathbb{E}[(Y - \hat{F}(X))^2]$ with $\hat{F}(x) = \hat{\mathbb{E}}[Y|X = x]$ by randomly splitting the data into 297 training and 33 test observations and averaging 50 times over such random partitions. Table 2 displays the results. We conclude that L_2Boost with component-wise splines is better than with trees and that it is among the best, together with classical backfitting of additive models. Moreover, estimation of the number of boosting iterations works very satisfactorily, exhibiting a performance which is very close to the optimum.

Next, we show a simulated example in very high dimensions relative to sample size, where L_2Boost as a stagewise method is better than backfitting for additive models with

method	mean squared error
L_2Boost with component-wise spline	17.78
L_2Boost with stumps	21.33
additive model (backfitted)	17.41
MARS	18.09
with optimal no. of iterations	
L_2Boost with component-wise spline	17.50 (5)
L_2Boost with stumps	20.96 (26)

Table 2: Cross-validated test set mean squared errors for ozone data. L_2Boost with estimated and optimal (minimizing cross-validated test set error) number of boosting iterations, given in parentheses. The component-wise spline is a cubic smoothing spline with df = 5.

variable selection. The simulation model is,

$$Y = \sum_{j=1}^{100} \left(1 + (-1)^j A_j X_j + B_j \sin(6X_j)\right) \sum_{j=1}^{50} (1 + X_j/50) + \varepsilon,$$

$A_1, \ldots, A_{100}$ i.i.d. Unif($[0.6, 1]$) and

$B_1, \ldots, B_{100}$ i.i.d. Unif($[0.8, 1.2]$), independent from the A_j's,

$X \sim$ Unif($[0, 1]^{100}$) where all components are i.i.d. $\sim$ Unif($[0, 1]$),

$\varepsilon \sim \mathcal{N}(0, 2)$. (12)

Samples of size $n = 200$ are generated from model (12): for each model realization (realized coefficients $A_1, \ldots A_{100}, B_1, \ldots, B_{50}$), we generate 200 i.i.d. pairs (Y_i, X_i) ($i = 1, \ldots, n = 200$).

We use the same methods as above. However, we use classical additive modeling with a forward variable selection (inclusion) strategy because $d = 100$ is very large compared to $n = 200$. We evaluate $\mathbb{E}[(\hat{F}(X) - \mathbb{E}[Y|X])^2]$ (X a new test observation) at the true conditional expectation which can be done in simulations. Already a stump appeared to be too strong for L_2Boost and we therefore used shrunken learners $\nu \mathcal{S}$ with $\nu = 0.5$ chosen ad-hoc; we also allow the non-boosting procedures to be shrunken with $\nu = 0.5$. Table 4 shows the average performance over 10 simulations from model (12). L_2Boost is the winner over additive models and MARS, and the component-wise spline is a better learner than stumps. Note that only the methods in the upper part of Table 3 are fully data-driven (the additive fits use the number of variables minimizing the true mean squared error). We believe that it is mainly in such high-dimensional situations where boosting has a clear advantage over other flexible nonparametric methods; and not so much in examples where the dimension d is "mid-range" or even small relative to sample size.

For simulated data, assessing significance for differences in performance is easily possible. Because all the different methods are used on the same data realizations, pairwise comparisons should be made. Table 4 displays the p-values of paired Wilcoxon tests. We conclude for this simulation model (12) that L_2Boost with componentwise smoothing spline is best and significantly outperforms the more classical methods such as MARS or additive models.

14

method	mean squared error
L_2Boost with shrunken component-wise spline	11.87
L_2Boost with shrunken stumps	12.76
MARS	25.19
shrunken MARS	15.05
with optimal no. of iterations or variables	
L_2Boost with shrunken component-wise spline	10.69 (228)
L_2Boost with shrunken stumps	12.54 (209)
additive model (backfitted and forward selection)	16.61 (1)
shrunken additive model (backfitted and forward selection)	14.44 (19)

Table 3: Mean squared errors for simulated data from (12) with $n = 200$. L_2Boost with estimated and optimal (minimizing MSE) number of boosting iterations, given in parentheses; additive model with optimal number of selected variables, given in parentheses. The component-wise spline is a cubic smoothing spline with df $= 5$; shrinkage factor is always $\nu = 0.5$.

	L_2Boost shr. stumps	shr. MARS	shr. additive model
L_2Boost shr. comp. spline	0.193	0.010	0.004
L_2Boost shr. stumps	–	0.010	0.014
shr. MARS	–	–	0.695

Table 4: p-values from paired two-sided Wilcoxon-tests for equal mean squared errors. Simulated data from (12) with $n = 200$ and 10 independent model realizations. p-values of two-sided tests are always in favor of the method in the row, i.e. sign of test-statistic always points towards favoring the method in the row. L_2Boost with estimated number of boosting iterations, specification of the methods as in Table 3 ("shr." abbreviates shrunken).

5 L_2Boosting for classification

The L_2Boost algorithm can also be used for classification, and it enjoys a computational simplicity in comparison with AdaBoost and LogitBoost.

5.1 Two-class problem

Consider a training sample

$$(Y_1, X_1), \ldots, (Y_n, X_n) \text{ i.i.d., } Y_i \in \{-1, 1\}, \ X_i \in \mathbb{R}^d. \tag{13}$$

L_2Boosting then yields an estimate $\hat{F}_m$ for the unknown function $\mathbb{E}[Y|X = x] = 2p(x) - 1$, where $p(x) = \mathbb{P}[Y = 1|X = x]$; the classification rule is given by (5). Note that also in two-class problems, the generic functional gradient descent repeatedly fits some learner $h(x, \theta)$ taking values in $\mathbb{R}$.

In addition to the L_2Boost algorithm, we propose a modification called "L_2Boost with constraints" (L_2WCBoost). It proceeds as L_2Boost, except that $\hat{F}_m(x)$ is constrained to be in $[-1, 1]$: this is natural since the target $F(x) = \mathbb{E}[Y|X = x] \in [-1, 1]$ is also in this range.

L_2**WCBoost algorithm**

Step 1. $\hat{F}_0(x) = h(x; \hat{\theta}_{Y,X})$ by least squares fitting. Set $m = 0$.

Step 2. Compute residuals $U_i = Y_i - \hat{F}_m(X_i)$ $(i = 1, \ldots, n)$. Then, fit $(U_1, X_1) \ldots, (U_n, X_n)$ by least squares

$$\hat{f}_{m+1}(x) = h(x, \hat{\theta}_{U,X}).$$

Update

$$\tilde{F}_{m+1}(\cdot) = \hat{F}_m(\cdot) + \hat{f}_{m+1}(\cdot),$$
$$\hat{F}_{m+1}(x) = \text{sign}(\tilde{F}_{m+1}(x)) \min\left(1, |\tilde{F}_{m+1}(x)|\right).$$

Note that if $|\tilde{F}_{m+1}(x)| \leq 1$, the updating step is as in the L_2Boost algorithm.

Step 3. Increase m by one and go back to Step 2.

5.1.1   Theory

For theoretical purposes, consider the model in (13) but with fixed, real-valued predictors:

$$Y_i \in \{-1, 1\} \text{ independent, } \mathbb{P}[Y_i = 1|x_i] = p(x_i), \ x_i \in \mathbb{R} \ (i = 1, \ldots, n). \tag{14}$$

Estimating $\mathbb{E}[Y|x_i] = 2p(x_i) - 1$ in classification can be seen as estimating the regression function in a heteroscedastic model,

$$Y_i = 2p(x_i) - 1 + \epsilon_i \ (i = 1, \ldots, n),$$

where ε_i are independent, mean zero variables, but with variance $4p(x_i)(1-p(x_i))$. Because the variances are bounded by 1, the arguments in the regression case can be modified to give the optimal rates of convergence results for estimating p.

Theorem 4. *(optimality of L2Boost for smoothing splines in classification). Consider the model in (14). Suppose $p \in \mathcal{W}_2^{(\nu)}$, see (9), and $\mathcal{S}$ is a smoothing spline linear learner $g_r(\lambda_0)$ of degree r, corresponding to a fixed smoothing parameter λ_0, see (10). If $\nu \geq r$, then there is an $m = m(n) = O(n^{2r/(2\nu+1)}) \to \infty$ such that $\hat{F}_{m(n)}$ achieves the optimal minimax rate $n^{-2\nu/(2\nu+1)}$ of the smoother function class $\mathcal{W}_2^{(\nu)}$ for estimating $2p(\cdot) - 1$ in terms of MSE as defined in (8).*

A proof is given in the Appendix. It is well-known that 2 times the L_1-norm bounds from above the difference between the generalization error of a plug-in classifier (expected 0-1 loss error for classifying a new observation) and the Bayes Risk (cf. Theorem 2.3 of Devroye et al., 1996). Furthermore, the L_1-norm is upper bounded by the L_2-norm. It follows from the above Theorem 4 that if the underlying p belongs to one of the smooth function class $\mathcal{W}_2^{(\nu)}$, L_2Boosting converges to the average Bayes risk (ABR)

$$ABR = n^{-1} \sum_{i=1}^{n} \mathbb{P}[sign(2p(x_i) - 1) \neq Y_i].$$

Moreover, the L_2-bound implies the following.

Corollary 2. *Under the assumptions and specifications in Theorem 4,*

$$n^{-1}\sum_{i=1}^{n}\mathbb{P}[sign(\hat{F}_{m(n)}(x_i)) \neq Y_i] - ABR = O(n^{-\nu/(2\nu+1)}).$$

Because the functions in $\mathcal{W}_2^{(\nu)}$ are bounded, using Hoeffding's inequality we can show that for both the terms in the above Corollary 2, replacing the average by an expectation with respect to a density for x would cause an error of order $o(n^{-\nu/(2\nu+1)})$ since $\nu/(2\nu + 1) < 1/2$. Hence the above result also holds if we replace the averaging by an expectation with respect to a randomly chosen x with a design density and replace the ABR by the corresponding Bayes risk.

For the generalization error with respect to a random x, Yang (1999) shows that the $n^{-\nu/(2\nu+1)}$ rate above is also minimax optimal for classification over the Lipschitz family $\mathcal{L}_2^{(\nu)}$

$$\mathcal{L}_2^{(\nu)} = \{p : \|p(x+h) - p(x)\|_2 < Ch^\nu, \ \|p\|_2 < C\}, \tag{15}$$

where $\|p\|_2 = (\int p^2(x)dx)^{1/2}$. Yang (1999) uses a hypercube subclass in $\mathcal{L}_2^{(\nu)}$ to prove the lower bound rate $n^{-\nu/(2\nu+1)}$. This hypercube subclass also belongs to our $\mathcal{W}_2^{(\nu)}$. Hence the rate $n^{-\nu/(2\nu+1)}$ also serves as a lower bound for our $\mathcal{W}_2^{(\nu)}$. Thus we have proved that the minimax optimal rate $n^{-\nu/(2\nu+1)}$ of convergence for the classification problem over the global smoothness class $\mathcal{W}_2^{(\nu)}$.

Marron (1983) gives a faster classification minimax rate of convergence $n^{-2\nu/(2\nu+1)}$ for a more restrictive global smoothness class and the same rate holds for us if we adopt that class. On the other hand, Mammen and Tsybakov (1999) consider different function classes which are locally constrained near the decision boundary and show that a faster than the parametric rate n^{-1} can even be achieved. The local constrained classes are more natural in the classification setting with the 0-1 loss since, as seen from our smoothed 0-1 loss expansion in Section 6.1, the actions happen near the decision boundary. These new rates are achieved by avoiding the plug-in classification rules via estimating p. Instead, empirical minimization of the 0-1 loss function over regularized classes of decision regions is used. However, computationally such a minimization could be very difficult. It remains open whether boosting can achieve these new optimal convergence rates in Mammen and Tsybakov (1999).

5.2 Multi-class problem

The multi-class problem has response variables $Y_i \in \{1, 2, \ldots, J\}$, taking values in a finite set of labels. An estimated classifier, under equal misclassification costs, is then given by

$$\hat{\mathcal{C}}(x) = \operatorname{argmax}_{j\in\{1,\ldots,J\}}\hat{\mathbb{P}}[Y = j|X = x].$$

The conditional probability estimates $\hat{p}_j(x) = \hat{\mathbb{P}}[Y = j|X = x]$ $(j = 1, \ldots, J)$ can be constructed from J different two-class problems, where each two-class problem encodes the events $\{Y = j\}$ and $\{Y \neq j\}$: this is also known as the "one against all" approach, cf. Allwein et al. (2001). Thus, the L_2Boost algorithm for multi-class problems works as follows:

Step 1. Compute $\hat{F}_m^{(j)}(\cdot)$ as an estimate of $p_j(\cdot)$ with L_2- or L_2WCBoost $(j = 1, \ldots, J)$ on the basis of binary response variables

$$Y_i^{(j)} = \begin{cases} 1 \text{ if } Y_i = j \\ -1 \text{ if } Y_i \neq j \end{cases}, \; i = 1, \ldots, n.$$

Step 2. Construct the classifier as

$$\hat{\mathcal{C}}_m(x) = \operatorname{argmax}_{j \in \{1, \ldots, J\}} \hat{F}_m^{(j)}(x).$$

The optimality results from Theorem 4 carries over to the multi-class case. Of course, other codings of a multi-class problem into multiple two-class problems can be done, cf. Allwein et al. (2001). The "one against all" scheme did work well in the cases we were looking at, see also section 7.2.

6 Understanding the 0-1 loss in two class problems

6.1 Generalization error via tapered moments of the margin

We consider here again the two-class problem as in (13). The performance is often measure by the generalization error

$$\mathbb{P}[sign(\hat{F}_m(X)) \neq Y] = \mathbb{P}[Y\hat{F}_m(X) < 0] = \mathbb{E}[\mathbf{1}_{[Y\hat{F}_m(X)<0]}]. \tag{16}$$

where $\mathbb{P}$ and $\mathbb{E}$ are over all the random variables in the training set (13) and the testing observation $(Y, X) \in \{-1, 1\} \times \mathbb{R}^d$, which is independent of the training set. Insights about the expected zero-one loss function in (16), the generalization error, can be gained by approximating it, for theoretical purposes, with a smoothed version

$$\mathbb{E}[C_\gamma(Y\hat{F}_m(X))],$$
$$C_\gamma(z) = (1 - \frac{\exp(z/\gamma)}{2})\mathbf{1}_{[z<0]} + \frac{\exp(-z/\gamma)}{2}\mathbf{1}_{[z\geq0]}, \; \gamma > 0.$$

The parameter γ controls the quality of approximation.

Proposition 4. *Assume that the distribution of $Z = Y\hat{F}_m(X)$ has a density $g(z)$ which is bounded for z in a neighborhood around zero. Then,*

$$|\mathbb{P}[Y\hat{F}_m(X) < 0] - \mathbb{E}[C_\gamma(Y\hat{F}_m(X))]| = O(\gamma \log(\gamma^{-1})) \; (\gamma \to 0).$$

A proof is given in the Appendix. Proposition 4 shows that the generalization error can be approximated by an expected cost function which is infinitely often differentiable.

Proposition 4 motivates to study generalization error through $\mathbb{E}[C_\gamma(Z)]$ with $Z = Y\hat{F}_m(X)$. Applying a Taylor series expansion of $C_\gamma(\cdot)$ around $Z^* = YF(X)$ we obtain

$$\mathbb{E}[C_\gamma(Z)] = \mathbb{E}[C_\gamma(Z^*)] + \sum_{k=1}^{\infty} \frac{1}{k!} \mathbb{E}[C_\gamma^{(k)}(Z^*)(Z - Z^*)^k]. \tag{17}$$

Thereby, the variables Z and Z^* denote the so-called estimated and true margin: a positive margin denotes a correct classification (and vice-versa) and the actual value of the margin

describes closeness to the classification boundary $\{x : F(x) = 0\}$. The derivatives of $C_\gamma(\cdot)$ are

$$C_\gamma^{(k)}(z) = \frac{1}{\gamma^k} \exp(\frac{-|z|}{\gamma})(-\mathbf{1}_{[z<0]} + (-1)^k \mathbf{1}_{[z\geq 0]}). \tag{18}$$

Using conditioning on the test observations (Y, X), the moments can be expressed as

$$\mathbb{E}[C_\gamma^{(k)}(Z^*)(Z - Z^*)^k] = \sum_{y\in\{-1,1\}} \int C_\gamma^{(k)}(yF(x))y^k b_k(x)\mathbb{P}[Y = y|X = x]dP_X(x),$$

$b_k(x) = \mathbb{E}[(\hat{F}_m(x) - F(x))^k]$, where expectation is over the training set in (13). (19)

Thereby, $P_X(\cdot)$ denotes the distribution of the predictors X.

From (17) and (19) we see that the smooth approximation to the generalization error of any procedure is approximately, in addition to the approximate Bayes risk $\mathbb{E}[C_\gamma Z^*]$, the sum of moments $b_k(x)$, tapered by $C_\gamma^{(k)}(yF(x))/k!$ which decays very quickly as $yF(x)$ moves away from zero. This exploits from a different view the known fact that only the behavior of $b_k(x)$ in the neighborhood of the classification boundary $\{x; \ F(x) = 0\}$ matters to the generalization error.

The first two terms in the approximation (17) are the tapered bias- and the tapered L_2-term, see (19) with $k = 1$ and 2. The higher order terms can be expanded as terms of interactions between the centered moments and the bias term (all tapered),

$$b_k(x) = \mathbb{E}[(\hat{F}_m(x) - F(x))^k] = \sum_{j=0}^{k} \binom{k}{j} b_1(x)^k \mathbb{E}[(\hat{F}_m(x) - \mathbb{E}[\hat{F}_m(x)])^{k-j}]. \tag{20}$$

This seemingly trivial approximation has three important consequences. The first is that bias (after tapering) as the first term in (17) and multiplicative terms in higher moments, see (20), plays a bigger role in (smoothed) 0-1 loss classification than in L_2-regression. Second, in the case of boosting, since all the (tapered) centered moment terms in (19) are bounded by expressions with exponentially diminishing increments as boosting iterations m get large (see Section 3, particularly Theorem 2) we gain resistance against overfitting. In view of the additional tapering, this resistance is even stronger than in regression, see also the rejoinder of Friedman et al. (2000) for a relevant illustration. The third consequence of the approximation in (17), (19) and (20) is to suggest why the previous attempts were not successful at decomposing the 0-1 prediction (generalization) error into additive bias and variance terms (cf. Geman et al. 1992; Breiman, 1998, and references therein). This, because except for the first two terms, all other important terms include the bias-term also in a multiplicative fashion (see (20)) for each term in the summation (17), instead of a pure additive way.

We conclude heuristically that the exponentially diminishing centered moment increase with the number of boosting iterations (as stated in Theorem 2), together with the tapering in the smoothed 0-1 loss yield the overall, often strong, overfitting-resistance performance of boosting in classification.

6.2 Acceleration of F and classification noise

As seen in Section 6.1, resistance against overfitting is closely related to the behavior of $F(\cdot)$ at the classification boundary. If the true $F(\cdot)$ moves away quickly from the

classification boundary $\{x; F(x) = 0\}$, the relevant tapering weights $C_\gamma^{(k)}(yF(x))$ decay very fast. This can be measured with $\text{grad}(F(x))|_{x=0}$, the gradient of F at zero. $F(\cdot)$ is said to have a large acceleration if its gradient is large (element-wise in absolute values, or in Euclidean norm). Thus, a large acceleration of $F(\cdot)$ should result in strong resistance against overfitting in boosting.

Noise negatively affects the acceleration of $F(\cdot)$. Noise in model (13), often called "classification noise", can be thought of in a constructive way. Consider a random variable $W \in \{-1, 1\}$, independent from (Y, X) with $\mathbb{P}[W = -1] = \pi$, $0 \leq \pi \leq 1/2$. The noisy response variable is

$$\tilde{Y} = WY, \tag{21}$$

changing the sign of Y with probability π. Its conditional probability is easily seen to be,

$$\mathbb{P}[\tilde{Y} = 1|X = x] = \mathbb{P}[Y = 1|X = x](1 - 2\pi) + \pi, \ 0 \leq \pi \leq 1/2. \tag{22}$$

Denote by $\tilde{F}(\cdot)$ the noisy version of $F(\cdot)$ with $\mathbb{P}[\tilde{Y} = 1|X = x]$ replacing $\mathbb{P}[Y = 1|X = x]$, either for $F(\cdot)$ being half of the log-odds ratio or the conditional expectation, see (3). A straightforward calculation then shows,

$$\text{grad}(\tilde{F}(x))|_{x=0} \searrow 0 \text{ as } \pi \nearrow 1/2. \tag{23}$$

The noisier the problem, the smaller the acceleration of $\tilde{F}(\cdot)$ and thus less resistance against overfitting since the tapering weights in (18) are becoming larger in noisy problems. This adds insights to the known empirical fact that boosting doesn't work well in noisy problems, see Dietterich (2000). Typically, overfitting then kicks in early and many learners are too strong in noisy problems. Using LogitBoost (Friedman et al., 2000), this effect is demonstrated in Figure 3 for the breast cancer data available at (http://www.ics.uci.edu/~mlearn/MLRepository) which has been analyzed by many others. We see there that already a stump becomes too strong for LogitBoost in the 25% or 40% noise added breast cancer data.

Of course, adding noise makes the classification problem harder. The optimal Bayes classifier $\tilde{\mathcal{C}}_{Bayes}(\cdot)$ in the noisy problem has generalization error

$$\mathbb{P}[\tilde{Y}(X) \neq \tilde{\mathcal{C}}_{Bayes}(X)] = \pi + (1 - 2\pi)\mathbb{P}[Y(X) \neq \mathcal{C}_{Bayes}(X)], \ 0 \leq \pi \leq 1/2,$$

relating it to the Bayes classifier $\mathcal{C}_{Bayes}$ of the original non-noisy problem. This is easily derived using (22). With high noise, the expression is largely dominated by the constant term π, indicating that there isn't much to gain by using a clever classifier (say close to optimal Bayes) instead of a naive one. Even more so, the *relative* improvement, which is often used to demonstrate the better performance of a powerful classifier (over a benchmark), becomes less dramatic due to the high noise level causing a large misclassification benchmark rate.

7 Comparing L_2Boost with LogitBoost on real and simulated data sets

We compare L_2Boost, our L_2WCBoost and LogitBoost using tree learners and our component-wise smoothing spline learner from Section 4.1.
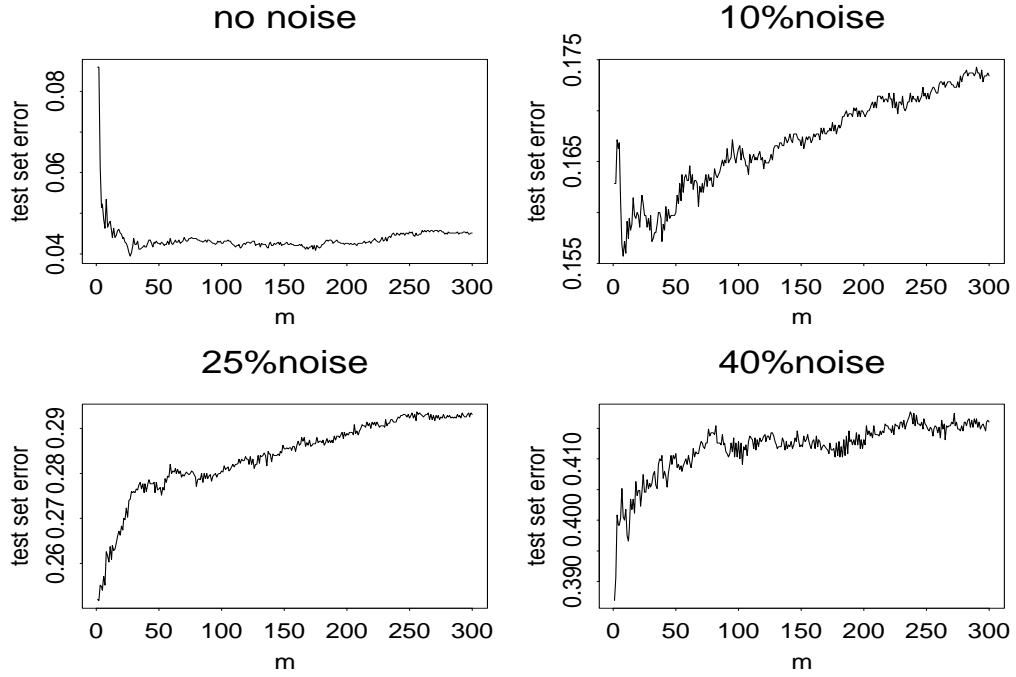
Figure 3: Test set misclassification errors for LogitBoost with stumps. Breast cancer data with different noise levels π (in %) as described in (21).

7.1 Two-class problems

We consider first a variety of 2-class problems from the UCI machine learning repository (http://www.ics.uci.edu/~mlearn/MLRepository): Breast cancer, Ionosphere, Monk 1 (full data set with $n = 432$) and the Heart, Sonar and Australian credit data sets from the Statlog project. Monk 1 has Bayes error equal to zero. The estimated test set misclassification errors, using an average of 50 random divisions into training with 90% and test set with 10% of the data, are given in Table 5. The comparison is made when using the optimal (with respect to cross-validated test set error) number of boosting iterations for every boosting algorithm; these numbers are given in parentheses. As has been well documented earlier (cf. Breiman (1998), Dietterich (2000), Friedman et al. (2000)), boosting with trees is usually much better than CART, displayed in Table 6. We use here two versions of CART: one where an initial large tree is pruned with the amount of pruning estimated by cross-validation, and the other using the default settings in S-Plus for an unpruned tree. See Table 6. The component-wise learner is used only for the breast cancer data with ordinal predictors and for the sonar data with continuous predictors. For the latter, spline smoothing makes most sense and we also found empirically that it is there where it improves upon the tree-based stumps. Estimating the number of boosting iterations can be done by 5-fold cross-validation as described in section 3.3. The effect of using the optimal number of boosting iterations instead of an estimated number is typically small: for example with the breast cancer data, estimating the number of boosting iterations by 5-fold cross-validation yields a final performance for L_2WCBoost with stumps as then 0.043 instead of 0.040 when using the optimal number of stopping. L_2WCBoost performs overall a bit better than L_2Boost, although in half of the data sets L_2Boost was better. LogitBoost was better than L_2Boost in 4 out of 6 and

dataset	n	d	learner	L_2Boost	L_2WCBoost	LogitBoost
Breast cancer	699	9	stumps	0.037 (176)	0.040 (275)	0.039 (27)
			comp. spline	0.036 (126)	0.043 (73)	0.038 (5)
Sonar	210	60	stumps	0.228 (62)	0.190 (335)	0.158 (228)
			comp. spline	0.178 (51)	0.168 (47)	0.148 (122)
Ionosphere	351	34	stumps	0.088 (25)	0.079 (123)	0.070 (157)
Heart (without costs)	270	13	stumps	0.167 (4)	0.175 (3)	0.159 (3)
Australian credit	690	14	stumps	0.123 (22)	0.123 (19)	0.131 (16)
Monk 1	432	7	larger tree	0.002 (42)	0.004 (12)	0.000 (6)

Table 5: Test set misclassification errors for L_2Boost, L_2WCBoost (with constraints) and LogitBoost. Optimal (minimizing cross-validated test set error) number of boosts is given in parentheses; if the optimum is not unique, the minimum is given. "Larger tree" denotes a tree learner such that the ancestor nodes of the terminal leaves contain at most 10 observations: resulting average tree size (integer-rounded) for L_2WCBoost is 12 terminal nodes.

	Breast cancer	Sonar	Ionosphere	Heart	Australian credit	Monk 1
unpruned CART	0.060	0.277	0.136	0.233	0.151	0.164
pruned CART	0.067	0.308	0.104	0.271	0.146	0.138

Table 6: Test set misclassification errors for CART. Unpruned version with defaults from S-Plus; pruning with cost-complexity parameter, estimated via minimal cross-validated deviance.

better than L_2WCBoost in 5 out of 6 data sets, but most often only by a small amount. The biggest difference in performance is for the Sonar data which has the most extreme ratio of dimension d to sample size n. But also for this data set, the difference is far from being significant since sample size is much too small.

Therefore, we consider next simulated data to compare the L_2WCBoost (the slightly better of the L_2procedures) with the LogitBoost algorithm. It allows a more accurate comparison by choosing large test sets. We generate data with two classes from the model in Friedman et al. (2000) with a non-additive decision boundary,

$$X \sim \mathcal{N}_{10}(0, I), \; Y|X = x \sim 2 \text{ Bernoulli}(p(x)) - 1,$$

$$\log(p(x)/(1 - p(x))) = 10 \sum_{j=1}^{10} x_j (1 + \sum_{k=1}^{6} (-1)^k x_k). \tag{24}$$

The (training) sample size is $n = 2000$. It is interesting to consider the performance on a single training and test data which resembles the situation in practice. For that purpose we choose a very large test set of size 100'000 so that variability of the test set error given the training data is very small. Additionally, we consider an average performance over 10 independent realizations from the model: here, we choose test set size as 10'000 which is still large compared to the training sample size $n = 2000$. The latter is the same set-up as in Friedman et al. (2000).

Figure 4 displays the results on a single data set. L_2WCBoost and LogitBoost perform about equally well. There is a slight advantage for L_2WCBoost with the larger regression tree learner having about 9 terminal nodes. It has been pointed out by Friedman et

stumps; single data set

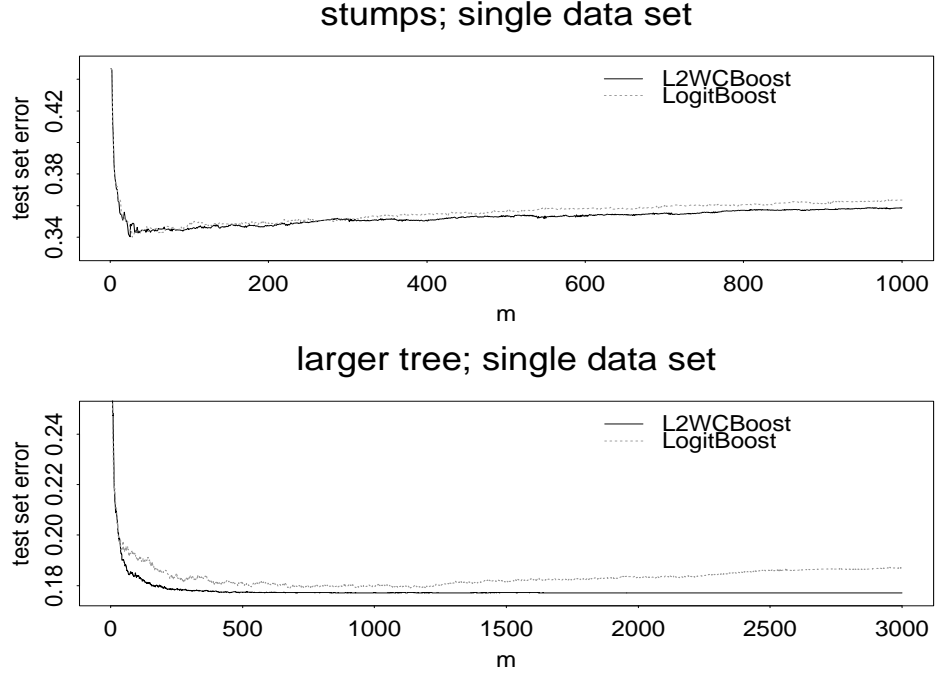## larger tree; single data set

Figure 4: Test set misclassification errors for a single realization from model (24). Top: stumps as learner. Bottom: larger tree as learner with at most 175 observations per terminal node (integer-rounded average tree size for L_2WCBoost is 9 terminal nodes).

larger tree; 10 simulations

Figure 5: Test set misclassification errors averaged over 10 realizations from model (24). Larger tree as learner with at most 175 observations per terminal node (integer-rounded average tree size for L_2WCBoost is 9 terminal nodes).

al. (2000) that stumps are not good learners because the true decision boundary is non-additive. With stumps as learners, the optimally (minimizing test set error) stopped LogitBoost has a tiny advantage over the optimally stopped L_2WCBoost (by about 0.13 estimated standard errors for the test set error estimate, given the training data). With larger trees, L_2Boost has a more substantial advantage over LogitBoost (by about 1.54 standard errors, given the data).

Figure 5 shows the averaged performance over 10 independent realizations with larger trees as learner: it indicates a more substantial advantage than on the single data represented by Figure 4. With 10 independent realizations, we test whether one of the optimally (minimizing test set error) stopped Boosting algorithms yields a significantly better test set misclassification error. In 9 out of the 10 simulations, optimally stopped L_2WCBoost was better than LogitBoost. The p-values for testing the hypothesis of equal performance against the two-sided alternative are given in Table 7. Thus, for the model (24) from

	t-test	Wilcoxon-test	sign-test
p-value	0.0015	0.0039	0.0215

Table 7: Comparison of L_2WCBoost and LogitBoost: two-sided testing for equal test set performance. Low p-values are always in favor of L_2WCBoost.

Friedman et al. (2000), we find a significant advantage of L_2WCBoost over LogitBoost.

It is not so surprising that L_2WCBoost and LogitBoost perform similarly. In nonparametric problems, loss functions are used locally in a neighborhood of a fitting point $x \in \mathbb{R}^d$; an example is the local likelihood framework, cf. Loader (1999). But locally, the L_2- and negative log-likelihood (with Bernoulli distribution) losses have the same minimizers.

7.2 Multi-class problems

We also run the L_2WCBoost algorithm using the "one-against all" approach described in section 7.2 on two often analyzed multi-class problems. The results are given in Table 8. The test set error curve remained essentially flat after the optimal number of boosting

dataset	n_{train}, n_{test}	d	no. classes	learner	error
Satimage	4435, 2000	36	6	stumps	0.110 (592)
				$\approx$ 8 terminal node tree	0.096 (148)
Letter	16000, 4000	16	26	$\approx$ 30 terminal node tree	0.029 (460)

Table 8: Test set misclassification errors for L_2WCBoost. Optimal number of boosts (minimizing test set error), is given in parentheses; if the optimum is not unique, the minimum is given. The approximate tree size is the integer-rounded average of tree sizes used during all 800 (satimage) or 1000 (letter) boosting iterations.

iterations: for the satimage data, the range of the test set error was $[0.096, 0.105]$ for iteration numbers 148 - 800 (pre-specified maximum); while for the letter data, the range was $[0.029, 0.031]$ for iteration numbers 460-1000 (pre-specified maximum). Therefore, a similar performance is expected with estimated number of iterations.

In comparison to our results with L_2WCBoost, Friedman et al. (2000) obtained the following with their LogitBoost algorithm using the fixed number of 200 boosting iterations: 0.102 with stumps and 0.088 with 8 terminal node tree for the satimage data; and 0.033 with 8 terminal node tree for the letter data. We have also tuned CART and

obtained the following test set error rates: 0.146 for satimage and 0.133 for the letter data set.

Thus, we find here a similar situation as for the 2-class problems: L_2WCBoost is about as good as LogitBoost and both of them are better than CART (by a significant amount for the letter data set).

8   Discussion and concluding remarks

As seen in Section 3, the computationally simple L_2Boosting is successful if the learner is sufficiently weak (sufficiently low variance), see also Figure 2. If the learner is too strong (too complex), then even at the first boosting iteration the MSE is not improved over the original learner. Also in the classification case it is likely that the generalization error will then increase with the number of boosting steps, stabilizing eventually due to the numerical convergence of the boosting algorithm to the fully saturated model (at least for linear learners with eigenvalues bounded away from zero). Of course, the weakness of a learner depends also on the underlying signal to noise ratio, see also assertion (2) in Theorem 1. But the degree of weakness of a learner can always be increased by an additional shrinkage using $\mathcal{S}_\nu = \nu \mathcal{S}$ with shrinkage factor $0 < \nu < 1$. For L_2Boost with one-dimensional predictors, we have presented asymptotic results in Section 3.2, stating that boosting weak smoothing splines is as good or even better than using an optimal smoothing spline since boosting adapts to unknown smoothness. In high dimensions, we see most practical advantages of boosting: we have exemplified this in Section 4.2.

The fact that the effectiveness of L_2Boost depends on the strength of the learner and the underlying problem is also (empirically) true for other variants of boosting. We show in Figure 6 some results for LogitBoost with trees and with projection pursuit learners having one ridge function (one term) for the breast cancer data. All tree learners are sufficiently weak for the problem and the shrunken large tree seems best. Interestingly, the projection pursuit learner is already strong and boosting does not pay off. This matches the intuition that projection pursuit is very flexible, complex and hence strong. Boosting projection pursuit a second time ($m = 2$) reduces performance: this estimate with $m = 2$ can then be improved by further boosting. Eventually the test set error curve exhibits overfitting and stabilizes at a relatively high value. We observed a similar pattern with projection pursuit learners in another simulated example. Such a multi-minimum phenomenon is typically not found with tree learners, but it is not inconsistent with our theoretical arguments.

Most of the empirical studies of boosting in the literature use a tree-structured learner. Their complexity is often low because they are themselves fitted by a stagewise selection of variables and split points, as in CART and other tree algorithms (while the complexity would be much higher, or it would be a much stronger learner, if the tree were fitted by a global search – which is of course infeasible). When the predictor space is high-dimensional, the substantial amount of variable selection done by a tree procedure is particularly helpful in leading to a low complexity (or weak) learner. And this feature may be very desirable.

In this paper, we mainly investigated the computationally simple L_2Boost by taking advantage of its analytical tractability, and we also demonstrated its practical effectiveness. To summarize, we showed that

1. L_2Boost is appropriate both for regression and classification. It leads to competitive

Figure 6: Test set misclassification errors of LogitBoost for breast cancer data. Top: decision tree learners, namely stumps (solid line), large unpruned tree (dotted line) and shrunken large unpruned tree (dashed line) with shrinkage factor $\nu = 0.01$. Bottom: projection pursuit (PPR) learners , namely one term PPR (solid line) and shrunken one term PPR (dotted line) with shrinkage factor $\nu = 0.01$.

performance, also in classification relative to LogitBoost.

2. L_2Boost is a stagewise fitting procedure with the iteration m acting as the smoothing or regularization parameter (this is also true with other boosting algorithms). In the linear learner case, m controls a new exponential bias-variance trade-off.

3. L_2Boost with smoothing splines results in optimal minimax rates of convergence, both in regression and classification. Moreover, the algorithm adapts to unknown smoothness.

4. Boosting learners which involve only one (selected) predictor variable yields an additive model fit. We propose component-wise cubic smoothing splines, which are of such type, and we believe that they are often better learners than tree-structured stumps, especially for continuous predictors.

5. Weighting observations is *not* used for L_2Boost: we doubt that the success of general boosting algorithms is due to "giving large weights to heavily misclassified instances" as Freund and Schapire (1996) conjectured for AdaBoost. Weighting in AdaBoost and LogitBoost comes as a consequence of the choice of the loss function, and is likely not the reason for their successes. It is interesting to note here Breiman's (2000) conjecture that even in the case of AdaBoost, weighting is not the reason for success.

26

6. A simple expansion of the smoothed 0-1 loss reveals new insights into the classification problem, particularly an additional resistance of boosting against overfitting.

Acknowledgments

Appendix

Proof of Proposition 1.
For L_2Boost with cost function $C(y, u) = (y - u)^2/2$, the negative gradient in stage j is the classical residual vector $u_j = Y - F_{j-1}$ and $F_j = F_{j-1} + f_j$ (there is no need for a line search) with $f_j = \mathcal{S}u_j$. Thus,

$$u_j = Y - F_{j-1} = u_{j-1} - \mathcal{S}u_{j-1} = (I - \mathcal{S})u_{j-1}, \; j = 1, 2, \ldots, m,$$

implying $u_j = (I - \mathcal{S})^j Y$ for $j = 1, 2, \ldots, m$. Since $F_0 = \mathcal{S}Y$ we obtain $\hat{F}_m = \sum_{j=0}^{m} \mathcal{S}(I - \mathcal{S})^j Y$. Using a telescope-sum argument, this equals $(I - (I - \mathcal{S})^{m+1})Y$. □

Proof of Proposition 3.
The bias term is

$$\text{bias}^2(m, \mathcal{S}; f) = (\mathbb{E}[\mathcal{B}_m Y] - f)^T (\mathbb{E}[\mathcal{B}_m Y] - f) = ((\mathcal{B}_m - I)f)^T ((\mathcal{B}_m - I)f).$$

According to (7), using orthonormality of U,

$$\mathcal{B}_m - I = U(D_m - I)U^T = U\text{diag}(-(1 - \lambda_k)^{m+1})U^T.$$

Thus, again by orthonormality of U, the formula for the bias follows.
For the variance, consider

$$\text{Cov}(\mathcal{B}_m Y) = \mathcal{B}_m \text{Cov}(Y)\mathcal{B}_m^T = \sigma^2 \mathcal{B}_m \mathcal{B}_m^T = \sigma^2 U\text{diag}((1 - (1 - \lambda_k)^{m+1})^2)U^T,$$

using (7) and orthonormality of U. Then,

$$\text{variance}(m, \mathcal{S}; \sigma^2) = tr[\text{Cov}(\mathcal{B}_m Y)] = \sigma^2 \sum_{k=1}^{n} (1 - (1 - \lambda_k)^{m+1})^2,$$

again using orthonormality of U. □

Proof of Theorem 1.
Assertion (1) is an immediate consequence of Proposition 3.
Without loss of generality, assume $\lambda_k < 1$ for all k. If not, restrict the summation to those k's which satisfy $\lambda_k < 1$.
(i) Denote by $\mu = U^T f \in \mathbb{R}^n$. For $x \geq 0$, let

$$g(x) = n^{-1} \sum_{k=1}^{n} \mu_k^2 (1 - \lambda_k)^{2x+2} + \sigma^2 n^{-1} \sum_{k=1}^{n} (1 - (1 - \lambda_k)^{x+1})^2.$$

Then

$$g(m) = MSE(m, \mathcal{S}; f, \sigma^2) = \text{bias}^2(m, \mathcal{S}; f) + \text{variance}(m, \mathcal{S}; \sigma^2).$$

It is easy to derive

$$g'(x) = 2n^{-1} \sum_{k=1}^{n} [(\mu_k^2 + \sigma^2)(1 - \lambda_k)^{x+1} - \sigma^2](1 - \lambda_k)^{x+1} \log(1 - \lambda_k).$$

It follows that

$$g'(0) = 2n^{-1} \sum_{k=1}^{n} [(\mu_k^2 + \sigma^2)(1 - \lambda_k) - \sigma^2](1 - \lambda_k) \log(1 - \lambda_k),$$

$$g'(1) = 2n^{-1} \sum_{k=1}^{n} [(\mu_k^2 + \sigma^2)(1 - \lambda_k)^2 - \sigma^2)](1 - \lambda_k)^2 \log(1 - \lambda_k)$$

which are both negative under the inequality condition in (i); also $g'(x) < 0$ for $x \in (0, 1)$ under this condition. Hence $g(1) < g(0)$ which means boosting improves.

(ii) Rewrite

$$g(x) = n^{-1} \sum_{k=1}^{n} [(\mu_k^2 + \sigma^2)(1 - \lambda_k)^{x+1} - 2\sigma^2](1 - \lambda_k)^{x+1} + \sigma^2.$$

Since for all k (with $\lambda_k < 1$), $(\mu_k^2 + \sigma^2)(1 - \lambda_k)^{x+1} - 2\sigma^2 \to -2\sigma^2$ as $x \to \infty$, there exists an m such that $(\mu_k^2 + \sigma^2)(1 - \lambda_k)^{m+1} - 2\sigma^2 \le -\sigma^2$ for all k (with $\lambda_k < 1$). It follows that

$$g(m) \le -n^{-1} \sum_{k=1}^{n} (1 - \lambda_k)^{m+1} \sigma^2 + \sigma^2 < \sigma^2.$$

It is obvious that $g(m) \to \sigma^2$ as $m \to \infty$. $\qquad\qquad\qquad \square$

Proof of Theorem 2.
Write the summands with the higher order moment as

$$\mathbb{E}[(\hat{F}_m(x_i) - f(x_i))^p] = \sum_{j=0}^{p} \binom{p}{j} b_m(x_i)^j \mathbb{E}[(\hat{F}_m(x_i) - \mathbb{E}[\hat{F}_m(x_i)])^{p-j}], \qquad (25)$$

where $b_m(x_i) = \mathbb{E}[\hat{F}_m(x_i) - f(x_i)]$ is the bias term. Thus, we have transformed the higher order moment as a sum of higher order *centered* moments with the bias as a multiplier which goes to zero as $m \to \infty$. The centered moments can be written as

$$(\hat{F}_m(x_i) - \mathbb{E}[\hat{F}_m(x_i)])^q = (\mathcal{B}_m Y - \mathbb{E}[\mathcal{B}_m Y])_i^q$$
$$= (\mathcal{B}_m \varepsilon)_i^q = ((I - (I - \mathcal{S})^{m+1})\varepsilon)_i^q = (\varepsilon - (I - \mathcal{S})^{m+1}\varepsilon)_i^q,$$

where we used assertion 1 from Proposition 1. Since $(I - \mathcal{S})^{m+1}$ is a map (matrix) taking values exponentially close to zero as $m \to \infty$, we obtain

$$\mathbb{E}[(\hat{F}_m(x_i) - \mathbb{E}[\hat{F}_m(x_i)])^q] = \mathbb{E}[\varepsilon_i^q] + O(\exp(-C_q m)) \ (m \to \infty)$$

for some constant $C_q > 0$. From this last bound and using (25) together with the fact that the bias $b_m(x_i) = O(\exp(-C_b m))$ $(m \to \infty)$ for some constant $C_b > 0$ (see Theorem 1), we complete the proof. $\qquad\square$

Proof of Theorem 3.

Let $\mathcal{S}$ be the smoothing spline operator corresponding to smoothness r and with smoothing parameter $c = \lambda_0$ (to avoid notational confusion with eigenvalues). It is well-known (cf. Utreras, 1983; Wahba, 1990, p.61) that the eigenvalues of $\mathcal{S}$ take the form in decreasing order

$$\lambda_1 = ... = \lambda_r = 1, \quad \lambda_k = \frac{n q_{k,n}}{n \lambda_0 + n q_{k,n}} \text{ for } k = r+1, ..., n.$$

Moreover, for n large, $q_{k,n} \approx A k^{-2r} := A q_k$ where A is universal and depends on the asymptotic density of the design points x_i. For the true function $f \in \mathcal{W}_2^{(\nu)}$,

$$\frac{1}{n} \sum_{k=r+1}^{n} \mu_k^2 k^{2\nu} \le M < \infty.$$

Let $c_0 = c/A$, then

$$\lambda_k \approx \frac{q_k}{c_0 + q_k} \text{ for } k = r+1, ..., n.$$

Then the bias term can be bounded as follows.

$$
\begin{aligned}
\text{bias}^2(m, \mathcal{S}; f) &= \frac{1}{n} \sum_{k=r+1}^{n} (1 - \lambda_k)^{2m+2} \mu_k^2 \\
&\approx \frac{1}{n} \sum_{k=r+1}^{n} (1 - q_k/(c_0 + q_k))^{2m+2} k^{-2\nu} \mu_k^2 k^{2\nu} \\
&\le \max_{k=r+1,...,n} (1 - q_k/(c_0 + q_k))^{2m+2} k^{-2\nu} \times \frac{1}{n} \sum_{k=r+1}^{n} \mu_k^2 k^{2\nu}. \\
&= \max_{k=r+1,...,n} \exp(h(k)) \times \frac{1}{n} \sum_{k=r+1}^{n} \mu_k^2 k^{2\nu},
\end{aligned}
$$

where

$$
\begin{aligned}
h(x) &= \log[(1 - x^{-2r}/(c_0 + x^{-2r}))^{2m+2} x^{-2\nu}] \\
&= (2m+2) \log(1 - 1/(c_0 x^{2r} + 1)) - 2\nu \log(x).
\end{aligned}
$$

Taking derivative gives

$$h'(x) = \frac{2r}{x} \frac{1}{c_0 x^{2r} + 1} [(2m+2) - \frac{\nu}{r}(c_0 x^{2r} + 1)].$$

Hence for any given positive integer n_1, if $x \le n_1$ and $m \ge \frac{\nu}{2r}(c_0 n_1^{2r} + 1) - 1$, $h(x)$ is increasing and so is $\exp(h(x))$, and

$$\exp(h(x)) \le \exp((h(n_1)) = (1 - 1/(c_0 n_1^{2r} + 1))^{2m+2} n_1^{-2\nu}.$$

On $[n_1 + 1, n]$,
$$\exp(h(x)) \leq (1 - 1/(c_0 n^{2r} + 1))^{2m+2} n_1^{-2\nu}.$$

Putting them together we get for any given n_1 and $m \geq \frac{\nu}{2r}(c_0 n_1^{2r} + 1) - 1$,
$$\text{bias}^2(m, \mathcal{S}; f) \leq M n_1^{-2\nu}[2(1 - 1/(c_0 n^{2r} + 1))^{2m+2}]$$

which is of the order $O(n_1^{-2\nu})$ for $n_1 \to \infty$ and $n_1 \leq n$.

Now let's deal with the variance term. For any $n_1 > r$,

$$
\begin{aligned}
\text{variance}(m, \mathcal{S}; \sigma^2) &= \frac{\sigma^2}{n}\{r + \sum_{k=r+1}^{n}[1 - (1 - \lambda_k)^{m+1}]^2 \\
&\leq \frac{\sigma^2 n_1}{n} + \frac{1}{n}\sum_{k=n_1+1}^{n}[1 - (1 - \lambda_k)^{m+1}]^2 := I_1 + I_2.
\end{aligned}
$$

Because $(1 - x)^a \geq 1 - ax$ for any $x \in [0, 1]$ and $a \geq 1$,
$$1 - (1 - \lambda_k)^{m+1} \leq 1 - [1 - (m + 1)\lambda_k] = (m + 1)\lambda_k.$$

It follows that

$$
\begin{aligned}
I_2 &\leq \frac{1}{n}\sum_{k=n_1+1}^{n}(m+1)^2\lambda_k^2 \approx \frac{(m+1)^2}{n}\sum_{k=n_1+1}^{n}\frac{1}{(c_0 k^{2r} + 2)^2} \\
&\leq \frac{(m+1)^2}{n}\sum_{k=n_1+1}^{n}\frac{1}{(c_0 k^{2r})^2} \leq \frac{(m+1)^2}{n}\int_{n_1}^{\infty}\frac{1}{(c_0 x^{2r})^2}dx \\
&= \frac{(m+1)^2}{c_0^2(4r-1)n}n_1/n_1^{4r} \leq O(n_1/n),
\end{aligned}
$$

if we take $m = m(n_1) = \frac{\nu}{2r}(c_0 n_1^{2r} + 1) - 1 = O(n_1^{2r})$. Hence for this choice of $m(n_1)$,
$$\text{variance}(m(n_1), \mathcal{S}; \sigma^2) \leq O(n_1/n).$$

Together with the bound for the bias we get
$$\frac{1}{n}MSE \leq O(n_1/n) + O(n_1^{-2\nu}),$$

which is minimized by taking $n_1 = O(n^{1/(2\nu+1)})$ and for $m(n) = m(n_1) = O(n^{2r/(2\nu+1)})$. The minimized MSE has the minimax optimal rate $O(n^{-2\nu/(2\nu+1)})$ of the smoother function class $\mathcal{W}_2^{(\nu)}$. $\qquad\square$

Proof of Proposition 4.
Denote by $C(z) = \mathbf{1}_{[z<0]}$. We first show that

$$\sup_{|z|>\gamma\log(\gamma^{-1})}|C(z) - C_\gamma(z)| = \gamma/2. \tag{26}$$

By symmetry, it suffices to consider $z > 0$,

$$\sup_{z>\gamma\log(\gamma^{-1})}|C(z) - C_\gamma(z)| = C_\gamma(\gamma\log(\gamma^{-1})) = \exp(-\log(\gamma^{-1}))/2 = \gamma/2,$$

30

proving (26).

On the other hand,

$$\int_{|z| \leq \gamma \log(\gamma^{-1})} |C(z) - C_\gamma(z)|g(z)dz \leq \sup_z |g(z)|\gamma \log(\gamma^{-1}) = O(\gamma \log(\gamma^{-1})). \qquad (27)$$

Hence, by (26) and (27),

$$
\begin{aligned}
&|\mathbb{E}[C(Z) - C_\gamma(Z)]| \\
\leq\ & (\int_{|z| > \gamma \log(\gamma^{-1})} + \int_{|z| \leq \gamma \log(\gamma^{-1})})|C(z) - C_\gamma(z)|g(z)dz \leq \gamma/2 + O(\gamma \log(\gamma^{-1})) \\
=\ & O(\gamma \log(\gamma^{-1}))\ (\gamma \to 0).
\end{aligned}
$$

$\square$

References

[1] Allwein, E., Schapire, R. and Singer, Y. (2001). Reducing multiclass to binary: a unifying approach for margin classifiers. J. Machine Learning Research **1**, 113–141.

[2] Buja, A. (2000). Comment on "Additive logistic regression: a statistical view of boosting". Ann. Statist. **28**, 387–391.

[3] Breiman, L. (1998). Arcing classifiers. Ann. Statist. **26**, 801-824.

[4] Breiman, L. (1999). Prediction games & arcing algorithms. Neural Computation **11**, 1493-1517.

[5] Breiman, L. (2000). Some infinity theory for predictor ensembles. Tech. Report 579, Dept. of Statist., Univ. of Calif., Berkeley.

[6] Devroye, L., Gyöye, L., and Lugosi, G. (1996). *A probabilistic theory of pattern recognition.* Springer, New York.

[7] Collins, M., Schapire, R.E. and Singer, Y. (2000). Logistic regression, AdaBoost and Bregman distances. Proc. Thirteenth Annual Conference Computational Learning Theory.

[8] Dietterich, T.G. (2000). An experimental comparison of three methods for constructing ensembles of decision trees: bagging, boosting, and randomization. Machine Learning **40**, 139–157.

[9] Freund, Y. (1995). Boosting a weak learning algorithm by majority. Information and Computation **121**, 256–285.

[10] Freund, Y. and Schapire, R.E. (1996). Experiments with a new boosting algorithm. In Machine Learning: Proc. Thirteenth International Conference, pp. 148–156. Morgan Kauffman, San Francisco.

[11] Friedman, J.H. (2001). Greedy function approximation: a gradient boosting machine. Ann. Statist. **29**, 1189–1232.

[12] Friedman, J.H., Hastie, T. and Tibshirani, R. (2000). Additive logistic regression: a statistical view of boosting. Ann. Statist. **28**, 337–407 (with discussion).

[13] Geman, S., Bienenstock, E. and Doursat, R. (1992). Neural networks and the bias/variance dilemma. Neural Computations **4**, 1–58.

[14] Gu, C. (1987). What happens when bootstrapping the smoothing spline? Commun. Statist. Part A - Theory Meth. **16**, 3275–3284.

[15] Hastie, T.J. and Tibshirani, R.J. (1990). *Generalized Additive Models*. Chapman & Hall.

[16] Jiang, W. (2000). Process consistency for AdaBoost. Tech. Report, Dept. of Statistics, Northwestern University.

[17] Loader, C. (1999). Local Regression and Likelihood. Series in Statistics and Computing. Springer, New York.

[18] Mammen, E. and Tsybakov, A.B. (1999). Smooth discriminant analysis. Ann. Statist. **27**, 1808–1829.

[19] Marron, J.S. (1983). Optimal rates of convergence to Bayes risk in nonparametric discrimination. Ann. Statist. **11**, 1142–1155.

[20] Mason, L., Baxter, J. Bartlett, P. and Frean, M. (1999). Functional gradient techniques for combining hypotheses. In *Advances in Large Margin Classifiers*. MIT Press.

[21] Schapire, R.E. (1990). The strength of weak learnability. Machine Learning **5**, 197–227.

[22] Schapire, R.E., Freund, Y., Bartlett, P. and Lee, W.S. (1998). Boosting the margin: A new explanation for the effectiveness of voting methods. Ann. Statist. **26**, 1651–1686.

[23] Tukey, J.W. (1977). *Exploratory data analysis*. Addison-Wesley, Reading, MA.

[24] Utreras, F. (1983). Natural spline functions, their associated eigenvalue problem. Numer. Math. **42**, 107–117.

[25] Wahba, G. (1990). *Spline Models for Observational Data*. Soc. for Industrial and Applied Mathematics.

[26] Yang, Y. (1999). Minimax nonparametric classification – Part I: rates of convergence. *IEEE Trans. Inform. Theory*, **45(7)**, 2271-2284.